

Package ‘rlsm’

November 4, 2017

Type Package

Title Least Squares Monte-Carlo

Version 1.0

Date 2016-07-10

Author Jeremy Yee

Maintainer Jeremy Yee <jeremyyee@outlook.com.au>

Description Least squares Monte Carlo and duality methods for Markov decision processes.

URL <https://github.com/YeeJeremy/rlsm>

License GPL

Depends StochasticProcess (>= 1.0)

Imports Rcpp (>= 0.11.6)

LinkingTo Rcpp, RcppArmadillo

NeedsCompilation yes

BugReports <https://github.com/YeeJeremy/rlsm/issues>

R topics documented:

AddDual	2
Bounds	3
FullTestPolicy	5
GetBounds	7
LSM	8
PathPolicy	10
TestPolicy	12
Index	14

AddDual

*AddDual***Description**

Compute additive duals.

Usage

```
AddDual(path, subsim, expected, Reward, Scrap, control, basis,
         basis_type, spline = FALSE, knots = matrix(NA, nrow = 1))
```

Arguments

path	3-D array representing sample paths. Entry [i,t] represents the state at time t for sample path i.
subsim	4-D array containing subsimulations. Entry [i,p,t] is for subsim i on path p at time t.
expected	3-D array representing the fitted coefficients for the continuation value function. Array [p,t] gives the fit for position p at time t.
Reward	User supplied function to represent the reward function. The function should take in the following arguments, in this order: <ul style="list-style-type: none"> • $n \times d$ matrix representing the n d-dimensional states. • A natural number representing the decision epoch. The function should output the following: <ul style="list-style-type: none"> • 3-D array with dimensions $n \times (a \times p)$ representing the rewards, where p is the number of positions and a is the number of actions in the problem. The $[i, a, p]$-th entry corresponds to the reward from applying the a-th action to the p-th position for the i-th state.
Scrap	User supplied function to represent the scrap function. The function should take in the following argument: <ul style="list-style-type: none"> • $n \times d$ matrix representing the n d-dimensional states.
control	Array representing the transition probabilities of the controlled Markov chain. Two possible inputs: <ul style="list-style-type: none"> • Matrix of dimension $n_pos \times n_action$, where entry [i,j] describes the next position after selecting action j at position i. • 3-D array with dimensions $n_pos \times n_action \times n_pos$, where entry [i,j,k] is the probability of moving to position k after applying action j to position i.
basis	Matrix specifying the regression basis. Zeros and ones.
intercept	Logical value indicating whether the intercept should be included.
basis_type	The type of basis functions to use: "power" and "laguerre".
spline	Logical value indicating whether linear splines should be used.
knots	Matrix indicating the location of the knots. If none, use NA.

Value

3-D array containing the additive duals. Entry $[i, j, t]$ is for path i and position j at time t .

Author(s)

Jeremy Yee

Examples

```
## Bermuda put option
step <- 0.02
mu <- 0.06 * step
vol <- 0.2 * sqrt(step)
n_dec <- 51
start <- 36
strike <- 40
## LSM
n_path <- 1000
path <- GBM(start, mu, vol, n_dec, n_path, TRUE)
control <- matrix(c(c(1, 1), c(2, 1)), nrow = 2, byrow = TRUE)
basis <- matrix(c(1, 1), nrow = 1)
knots <- matrix(c(30, 40, 50), nrow = 1)
Scrap <- function(state) {
  output <- matrix(data = 0, nrow = nrow(state), ncol = 2)
  output[, 2] <- exp(-mu * (n_dec - 1)) * pmax(strike - state, 0)
  return(output)
}
Reward <- function(state, time) {
  output <- array(data = 0, dim = c(nrow(state), 2, 2))
  output[, 2, 2] <- exp(-mu * (time - 1)) * pmax(strike - state, 0)
  return(output)
}
lsm <- LSM(path, Reward, Scrap, control, basis, TRUE, "power", TRUE, knots)
n_path2 <- 100
path2 <- GBM(start, mu, vol, n_dec, n_path2, TRUE)
policy <- PathPolicy(path2, lsm$expected, Reward, control, basis,
"power", TRUE, knots)
n_subsim <- 100
subsim <- NestedGBM(path2, mu, vol, n_subsim, TRUE)
mart <- AddDual(path2, subsim, lsm$expected, Reward, Scrap, control, basis, "power", TRUE, knots)
```

Bounds

Bounds

Description

Compute bound estimates using additive duals.

Usage

Bounds(path, Reward, Scrap, control, mart, path_action)

Arguments

path	3-D array representing sample paths. Entry $[i,t]$ represents the state at time t for sample path i .
Reward	<p>User supplied function to represent the reward function. The function should take in the following arguments, in this order:</p> <ul style="list-style-type: none"> • $n \times d$ matrix representing the n d-dimensional states. • A natural number representing the decision epoch. <p>The function should output the following:</p> <ul style="list-style-type: none"> • 3-D array with dimensions $n \times (a \times p)$ representing the rewards, where p is the number of positions and a is the number of actions in the problem. The $[i, a, p]$-th entry corresponds to the reward from applying the a-th action to the p-th position for the i-th state.
Scrap	<p>User supplied function to represent the scrap function. The function should take in the following argument:</p> <ul style="list-style-type: none"> • $n \times d$ matrix representing the n d-dimensional states.
control	<p>Array representing the transition probabilities of the controlled Markov chain. Two possible inputs:</p> <ul style="list-style-type: none"> • Matrix of dimension $n_pos \times n_action$, where entry $[i,j]$ describes the next position after selecting action j at position i. • 3-D array with dimensions $n_pos \times n_action \times n_pos$, where entry $[i,j,k]$ is the probability of moving to position k after applying action j to position i.
mart	3-D array containing the additive duals. Entry $[i, j, t]$ is for path i and position j at time t .
path_action	3-D array containing the prescribed policy. Entry $[i,p,t]$ is for path i and position p at time t .

Value

primal	3-D array containing the lower bound estimates. Entry $[i,p,t]$ is for path i and position p at time t .
dual	3-D array containing the lower bound estimates. Entry $[i,p,t]$ is for path i and position p at time t .

Author(s)

Jeremy Yee

Examples

```
## Bermuda put option
step <- 0.02
mu <- 0.06 * step
vol <- 0.2 * sqrt(step)
n_dec <- 51
start <- 36
strike <- 40
## LSM
n_path <- 1000
path <- GBM(start, mu, vol, n_dec, n_path, TRUE)
control <- matrix(c(c(1, 1), c(2, 1)), nrow = 2, byrow = TRUE)
basis <- matrix(c(1, 1), nrow = 1)
knots <- matrix(c(30, 40, 50), nrow = 1)
Scrap <- function(state) {
  output <- matrix(data = 0, nrow = nrow(state), ncol = 2)
  output[, 2] <- exp(-mu * (n_dec - 1)) * pmax(strike - state, 0)
  return(output)
}
Reward <- function(state, time) {
  output <- array(data = 0, dim = c(nrow(state), 2, 2))
  output[, 2, 2] <- exp(-mu * (time - 1)) * pmax(strike - state, 0)
  return(output)
}
lsm <- LSM(path, Reward, Scrap, control, basis, TRUE, "power", TRUE, knots)
n_path2 <- 100
path2 <- GBM(start, mu, vol, n_dec, n_path2, TRUE)
policy <- PathPolicy(path2, lsm$expected, Reward, control, basis,
  "power", TRUE, knots)
n_subsim <- 100
subsim <- NestedGBM(path2, mu, vol, n_subsim, TRUE)
mart <- AddDual(path2, subsim, lsm$expected, Reward, Scrap, control,
  basis, "power", TRUE, knots)
bounds <- Bounds(path2, Reward, Scrap, control, mart, policy)
```

FullTestPolicy

FullTestPolicy

Description

Full testing of prescribed policy for sample paths.

Usage

```
FullTestPolicy(start_position, path, control, Reward, Scrap, path_action)
```

Arguments

`start_position` Starting position.

path	3-D array representing sample paths. Entry $[i,t]$ represents the state at time t for sample path i .
control	<p>Array representing the transition probabilities of the controlled Markov chain. Two possible inputs:</p> <ul style="list-style-type: none"> Matrix of dimension $n_pos \times n_action$, where entry $[i,j]$ describes the next position after selecting action j at position i. 3-D array with dimensions $n_pos \times n_action \times n_pos$, where entry $[i,j,k]$ is the probability of moving to position k after applying action j to position i.
Reward	<p>User supplied function to represent the reward function. The function should take in the following arguments, in this order:</p> <ul style="list-style-type: none"> $n \times d$ matrix representing the n d-dimensional states. A natural number representing the decision epoch. <p>The function should output the following:</p> <ul style="list-style-type: none"> 3-D array with dimensions $n \times (a \times p)$ representing the rewards, where p is the number of positions and a is the number of actions in the problem. The $[i, a, p]$-th entry corresponds to the reward from applying the a-th action to the p-th position for the i-th state.
Scrap	<p>User supplied function to represent the scrap function. The function should take in the following argument:</p> <ul style="list-style-type: none"> $n \times d$ matrix representing the n d-dimensional states.
path_action	3-D array containing the prescribed policy. Entry $[i,p,t]$ is for path i and position p at time t .

Value

value	Array containing the path values.
position	Matrix containing the evolution of the position. Entry $[i,t]$ refers to the position at time t for sample path i .
action	Matrix containing the actions taken. Entry $[i,t]$ refers to the action at time t for sample path i .

Author(s)

Jeremy Yee

Examples

```
## Bermuda put option
step <- 0.02
mu <- 0.06 * step
vol <- 0.2 * sqrt(step)
n_dec <- 51
start <- 36
strike <- 40
## LSM
```

```

n_path <- 1000
path <- GBM(start, mu, vol, n_dec, n_path, TRUE)
control <- matrix(c(c(1, 1), c(2, 1)), nrow = 2, byrow = TRUE)
basis <- matrix(c(1, 1), nrow = 1)
knots <- matrix(c(30, 40, 50), nrow = 1)
Scrap <- function(state) {
  output <- matrix(data = 0, nrow = nrow(state), ncol = 2)
  output[, 2] <- exp(-mu * (n_dec - 1)) * pmax(strike - state, 0)
  return(output)
}
Reward <- function(state, time) {
  output <- array(data = 0, dim = c(nrow(state), 2, 2))
  output[, 2, 2] <- exp(-mu * (time - 1)) * pmax(strike - state, 0)
  return(output)
}
lsm <- LSM(path, Reward, Scrap, control, basis, TRUE, "power", TRUE, knots)
n_path2 <- 1000
path2 <- GBM(start, mu, vol, n_dec, n_path2, TRUE)
policy <- PathPolicy(path2, lsm$expected, Reward, control, basis,
"power", TRUE, knots)
test <- FullTestPolicy(2, path, control, Reward, Scrap, policy)

```

GetBounds

Confidence Bounds

Description

Confidence bounds for the value.

Usage

```
GetBounds(duality, alpha, position)
```

Arguments

duality	Object returned by the Bounds function.
alpha	Specifies the (1-alpha) confidence bounds.
position	Natural number indicating the starting position.

Value

Array representing the (1-alpha) confidence bounds for the value of the specified position.

Author(s)

Jeremy Yee

Examples

```
## Bermuda put option
step <- 0.02
mu <- 0.06 * step
vol <- 0.2 * sqrt(step)
n_dec <- 51
start <- 36
strike <- 40
## LSM
n_path <- 1000
path <- GBM(start, mu, vol, n_dec, n_path, TRUE)
control <- matrix(c(c(1, 1), c(2, 1)), nrow = 2, byrow = TRUE)
basis <- matrix(c(1, 1), nrow = 1)
knots <- matrix(c(30, 40, 50), nrow = 1)
Scrap <- function(state) {
  output <- matrix(data = 0, nrow = nrow(state), ncol = 2)
  output[, 2] <- exp(-mu * (n_dec - 1)) * pmax(strike - state, 0)
  return(output)
}
Reward <- function(state, time) {
  output <- array(data = 0, dim = c(nrow(state), 2, 2))
  output[, 2, 2] <- exp(-mu * (time - 1)) * pmax(strike - state, 0)
  return(output)
}
lsm <- LSM(path, Reward, Scrap, control, basis, TRUE, "power", TRUE, knots)
n_path2 <- 100
path2 <- GBM(start, mu, vol, n_dec, n_path2, TRUE)
policy <- PathPolicy(path2, lsm$expected, Reward, control, basis,
  "power", TRUE, knots)
n_subsim <- 100
subsim <- NestedGBM(path2, mu, vol, n_subsim, TRUE)
mart <- AddDual(path2, subsim, lsm$expected, Reward, Scrap, control,
  basis, "power", TRUE, knots)
bounds <- Bounds(path2, Reward, Scrap, control, mart, policy)
confidenceInterval <- GetBounds(bounds, 0.05, 2)
```

LSM

Least squares Monte Carlo

Description

Perform the least squares Monte Carlo algorithm.

Usage

```
LSM(path, Reward, Scrap, control, basis, intercept, basis_type,
  spline = FALSE, knots = matrix(NA, nrow = 1))
```


Arguments

path	3-D array representing sample paths. Entry $[i,j]$ represents the state at time j for sample path i .
Reward	User supplied function to represent the reward function. The function should take in the following arguments, in this order: <ul style="list-style-type: none"> • $n \times d$ matrix representing the n d-dimensional states. • A natural number representing the decision epoch. The function should output the following: <ul style="list-style-type: none"> • 3-D array with dimensions $n \times (a \times p)$ representing the rewards, where p is the number of positions and a is the number of actions in the problem. The $[i, a, p]$-th entry corresponds to the reward from applying the a-th action to the p-th position for the i-th state.
Scrap	User supplied function to represent the scrap function. The function should take in the following argument: <ul style="list-style-type: none"> • $n \times d$ matrix representing the n d-dimensional states.
control	Array representing the transition probabilities of the controlled Markov chain. Two possible inputs: <ul style="list-style-type: none"> • Matrix of dimension $n_pos \times n_action$, where entry $[i,j]$ describes the next position after selecting action j at position i. • 3-D array with dimensions $n_pos \times n_action \times n_pos$, where entry $[i,j,k]$ is the probability of moving to position k after applying action j to position i.
basis	Matrix specifying the regression basis. Zeros and ones.
intercept	Logical value indicating whether the intercept should be included.
basis_type	The type of basis functions to use: "power" and "laguerre".
spline	Logical value indicating whether linear splines should be used.
knots	Matrix indicating the location of the knots. If none, use NA.

Value

value	3-D array containing the path values. Entry $[i,p,t]$ is for path i and position p at time t .
expected	3-D array representing the fitted coefficients for the continuation value function. Array $[,p,t]$ gives the fit for position p at time t .

Author(s)

Jeremy Yee

Examples

```
## Bermuda put option
step <- 0.02
mu <- 0.06 * step
```

```

vol <- 0.2 * sqrt(step)
n_dec <- 51
start <- 36
strike <- 40
## LSM
n_path <- 1000
path <- GBM(start, mu, vol, n_dec, n_path, TRUE)
control <- matrix(c(c(1, 1), c(2, 1)), nrow = 2, byrow = TRUE)
basis <- matrix(c(1, 1), nrow = 1)
knots <- matrix(c(30, 40, 50), nrow = 1)
Scrap <- function(state) {
  output <- matrix(data = 0, nrow = nrow(state), ncol = 2)
  output[, 2] <- exp(-mu * (n_dec - 1)) * pmax(strike - state, 0)
  return(output)
}
Reward <- function(state, time) {
  output <- array(data = 0, dim = c(nrow(state), 2, 2))
  output[, 2, 2] <- exp(-mu * (time - 1)) * pmax(strike - state, 0)
  return(output)
}
lsm <- LSM(path, Reward, Scrap, control, basis, TRUE, "power", TRUE, knots)

```

PathPolicy

PathPolicy

Description

Obtaining the prescribed policy for sample paths

Usage

```
PathPolicy(path, expected, Reward, control, basis, intercept, basis_type,
  spline = FALSE, knots = matrix(NA, nrow = 1))
```

Arguments

- | | |
|----------|---|
| path | 3-D array representing sample paths. Entry [i,t] represents the state at time t for sample path i. |
| expected | 3-D array representing the fitted coefficients for the continuation value function. Array [p,t] gives the fit for position p at time t. |
| Reward | User supplied function to represent the reward function. The function should take in the following arguments, in this order: <ul style="list-style-type: none"> • $n \times d$ matrix representing the n d-dimensional states. • A natural number representing the decision epoch. |

The function should output the following:

	<ul style="list-style-type: none"> • 3-D array with dimensions $n \times (a \times p)$ representing the rewards, where p is the number of positions and a is the number of actions in the problem. The $[i, a, p]$-th entry corresponds to the reward from applying the a-th action to the p-th position for the i-th state.
control	<p>Array representing the transition probabilities of the controlled Markov chain. Two possible inputs:</p> <ul style="list-style-type: none"> • Matrix of dimension $n_pos \times n_action$, where entry $[i,j]$ describes the next position after selecting action j at position i. • 3-D array with dimensions $n_pos \times n_action \times n_pos$, where entry $[i,j,k]$ is the probability of moving to position k after applying action j to position i.
basis	Matrix specifying the regression basis. Zeros and ones.
intercept	Logical value indicating whether the intercept should be included.
basis_type	The type of basis functions to use: "power" and "laguerre".
spline	Logical value indicating whether linear splines should be used.
knots	Matrix indicating the location of the knots. If none, use NA.

Value

3-D array containing the prescribed policy. Entry $[i,p,t]$ is for path i and position p at time t .

Author(s)

Jeremy Yee

Examples

```
## Bermuda put option
step <- 0.02
mu <- 0.06 * step
vol <- 0.2 * sqrt(step)
n_dec <- 51
start <- 36
strike <- 40
## LSM
n_path <- 1000
path <- GBM(start, mu, vol, n_dec, n_path, TRUE)
control <- matrix(c(c(1, 1), c(2, 1)), nrow = 2, byrow = TRUE)
basis <- matrix(c(1, 1), nrow = 1)
knots <- matrix(c(30, 40, 50), nrow = 1)
Scrap <- function(state) {
  output <- matrix(data = 0, nrow = nrow(state), ncol = 2)
  output[, 2] <- exp(-mu * (n_dec - 1)) * pmax(strike - state, 0)
  return(output)
}
Reward <- function(state, time) {
  output <- array(data = 0, dim = c(nrow(state), 2, 2))
  output[, 2, 2] <- exp(-mu * (time - 1)) * pmax(strike - state, 0)
  return(output)
}
```

```

}
lsm <- LSM(path, Reward, Scrap, control, basis, TRUE, "power", TRUE, knots)
n_path2 <- 1000
path2 <- GBM(start, mu, vol, n_dec, n_path2, TRUE)
policy <- PathPolicy(path2, lsm$expected, Reward, control, basis, "power", TRUE, knots)

```

TestPolicy

TestPolicy

Description

Testing prescribed policy for sample paths.

Usage

```
TestPolicy(start_position, path, control, Reward, Scrap, path_action)
```

Arguments

- | | |
|----------------|--|
| start_position | Starting position. |
| path | 3-D array representing sample paths. Entry [i,t] represents the state at time t for sample path i. |
| control | <p>Array representing the transition probabilities of the controlled Markov chain. Two possible inputs:</p> <ul style="list-style-type: none"> • Matrix of dimension $n_{\text{pos}} \times n_{\text{action}}$, where entry [i,j] describes the next position after selecting action j at position i. • 3-D array with dimensions $n_{\text{pos}} \times n_{\text{action}} \times n_{\text{pos}}$, where entry [i,j,k] is the probability of moving to position k after applying action j to position i. |
| Reward | <p>User supplied function to represent the reward function. The function should take in the following arguments, in this order:</p> <ul style="list-style-type: none"> • $n \times d$ matrix representing the n d-dimensional states. • A natural number representing the decision epoch. <p>The function should output the following:</p> <ul style="list-style-type: none"> • 3-D array with dimensions $n \times (a \times p)$ representing the rewards, where p is the number of positions and a is the number of actions in the problem. The $[i, a, p]$-th entry corresponds to the reward from applying the a-th action to the p-th position for the i-th state. |
| Scrap | <p>User supplied function to represent the scrap function. The function should take in the following argument:</p> <ul style="list-style-type: none"> • $n \times d$ matrix representing the n d-dimensional states. |
| path_action | 3-D array containing the prescribed policy. Entry [i,p,t] is for path i and position p at time t. |

Value

Array containing the values for each path.

Author(s)

Jeremy Yee

Examples

```
## Bermuda put option
step <- 0.02
mu <- 0.06 * step
vol <- 0.2 * sqrt(step)
n_dec <- 51
start <- 36
strike <- 40
## LSM
n_path <- 1000
path <- GBM(start, mu, vol, n_dec, n_path, TRUE)
control <- matrix(c(c(1, 1), c(2, 1)), nrow = 2, byrow = TRUE)
basis <- matrix(c(1, 1), nrow = 1)
knots <- matrix(c(30, 40, 50), nrow = 1)
Scrap <- function(state) {
  output <- matrix(data = 0, nrow = nrow(state), ncol = 2)
  output[, 2] <- exp(-mu * (n_dec - 1)) * pmax(strike - state, 0)
  return(output)
}
Reward <- function(state, time) {
  output <- array(data = 0, dim = c(nrow(state), 2, 2))
  output[, 2, 2] <- exp(-mu * (time - 1)) * pmax(strike - state, 0)
  return(output)
}
lsm <- LSM(path, Reward, Scrap, control, basis, TRUE, "power", TRUE, knots)
n_path2 <- 1000
path2 <- GBM(start, mu, vol, n_dec, n_path2, TRUE)
policy <- PathPolicy(path2, lsm$expected, Reward, control, basis,
  "power", TRUE, knots)
test <- TestPolicy(2, path, control, Reward, Scrap, policy)
```

Index

AddDual, [2](#)

Bounds, [3](#)

FullTestPolicy, [5](#)

GetBounds, [7](#)

LSM, [8](#)

PathPolicy, [10](#)

TestPolicy, [12](#)